Olivetti face

pixels: the convex blend is *the output — you see a blur*

stored A   stored B   blurry ghost

No quantizer: the blend itself is the image — not a valid new face. So the clean output is a **verbatim stored image** (**cosine 1.00**).

RECALL

Protein sequence

a per-position $\arg\max$ quantizes *the blend — you get a recombinant*

SA-generated	L F V G N L P P D T T E E E L K E F F S Q F G P V V S V
nearest natural	L F I G G L A P Y T T E E N L K L F Y G Q W G K V V D V

teal *= sites recombined away from the nearest natural (each is still a residue the family uses there)*

every residue is one the family uses at that site (**71/71**, valid) • **38%** of sites recombined (62% identity to nearest, novel) • passes RRM **HMM**

The convex blend lives in embedding space; $\arg\max$ snaps each of the 71 sites to one residue, so the *average is never seen* — different sites inherit from different neighbors.

NOVEL + VALID